

▼ Keras -- MLPs on MNIST

```

# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()

# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()

print("Number of training examples : ", X_train.shape[0], "and each image is of shape (%d, %d)%(X_trai
print("Number of training examples : ", X_test.shape[0], "and each image is of shape (%d, %d)%(X_test.

    □→ Number of training examples : 60000 and each image is of shape (28, 28)
    Number of training examples : 10000 and each image is of shape (28, 28)

# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])

# after converting the input images from 3d to 2d vectors

print("Number of training examples : ", X_train.shape[0], "and each image is of shape (%d)%(X_train.sh
print("Number of training examples : ", X_test.shape[0], "and each image is of shape (%d)%(X_test.shap

    □→ Number of training examples : 60000 and each image is of shape (784)
    Number of training examples : 10000 and each image is of shape (784)

# An example data point
print(X_train[0])

```



```
# if we observe the above matrix each cell is having a value between 0-255  
# before we move to apply machine learning algorithms lets try to normalize the data  
#  $X \Rightarrow (X - X_{\min}) / (X_{\max} - X_{\min}) = X / 255$ 
```

```
X_train = X_train/255  
X_test = X_test/255
```

```
# example data point after normalizing  
print(X_train[0])
```




```

0.      0.      0.      0.      0.      0.04313725
0.74509804 0.99215686 0.2745098 0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.1372549 0.94509804
0.88235294 0.62745098 0.42352941 0.00392157 0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.31764706 0.94117647 0.99215686
0.99215686 0.46666667 0.09803922 0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.17647059 0.72941176 0.99215686 0.99215686
0.58823529 0.10588235 0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.0627451 0.36470588 0.98823529 0.99215686 0.73333333
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.

```

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])
```

```
# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs
```

```
Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)
```

```
print("After converting the output into a vector : ",Y_train[0])
```

→ Class label of first image : 5
 After converting the output into a vector : [0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]

Softmax classifier

```
# https://keras.io/getting-started/sequential-model-guide/
# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:
```

```
# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])
```

```
# You can also simply add layers via the .add() method:
```

```
# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))
```

```
###
```

```

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

#####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument suppo
# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

```
# some model parameters
```

```

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 30

```

MLP + BN + Dropout + AdamOptimizer + ReLu activation

Architecture 1: 784-366-100-10;
dropout: .5, .5

std = $\sqrt{2/\text{fan-in}}$ as it is relu

```

from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

model = Sequential()

model.add(Dense(366, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(100, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.14142, seed
model.add(BatchNormalization())

```

```
model.add(Dropout(0.5))

model.add(Dense(output_dim, activation='softmax'))

model.summary()

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_da
```

□→

Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_44 (Dense)	(None, 366)	287310
batch_normalization_31 (Batch Normalization)	(None, 366)	1464
dropout_31 (Dropout)	(None, 366)	0
dense_45 (Dense)	(None, 100)	36700
batch_normalization_32 (Batch Normalization)	(None, 100)	400
dropout_32 (Dropout)	(None, 100)	0
dense_46 (Dense)	(None, 10)	1010

Total params: 326,884
 Trainable params: 325,952
 Non-trainable params: 932

Train on 60000 samples, validate on 10000 samples

Epoch 1/30
 60000/60000 [=====] - 9s 158us/step - loss: 0.5750 - acc: 0.
 Epoch 2/30
 60000/60000 [=====] - 5s 88us/step - loss: 0.2700 - acc: 0.
 Epoch 3/30
 60000/60000 [=====] - 5s 91us/step - loss: 0.2075 - acc: 0.
 Epoch 4/30
 60000/60000 [=====] - 6s 94us/step - loss: 0.1752 - acc: 0.
 Epoch 5/30
 60000/60000 [=====] - 6s 95us/step - loss: 0.1544 - acc: 0.
 Epoch 6/30
 60000/60000 [=====] - 5s 87us/step - loss: 0.1380 - acc: 0.
 Epoch 7/30
 60000/60000 [=====] - 6s 93us/step - loss: 0.1268 - acc: 0.
 Epoch 8/30
 60000/60000 [=====] - 5s 90us/step - loss: 0.1162 - acc: 0.
 Epoch 9/30
 60000/60000 [=====] - 5s 90us/step - loss: 0.1100 - acc: 0.
 Epoch 10/30
 60000/60000 [=====] - 5s 89us/step - loss: 0.1020 - acc: 0.
 Epoch 11/30
 60000/60000 [=====] - 5s 91us/step - loss: 0.0961 - acc: 0.
 Epoch 12/30
 60000/60000 [=====] - 5s 91us/step - loss: 0.0918 - acc: 0.
 Epoch 13/30
 60000/60000 [=====] - 5s 90us/step - loss: 0.0866 - acc: 0.
 Epoch 14/30
 60000/60000 [=====] - 5s 92us/step - loss: 0.0821 - acc: 0.
 Epoch 15/30
 60000/60000 [=====] - 6s 92us/step - loss: 0.0762 - acc: 0.
 Epoch 16/30
 60000/60000 [=====] - 6s 92us/step - loss: 0.0750 - acc: 0.
 Epoch 17/30
 60000/60000 [=====] - 6s 92us/step - loss: 0.0722 - acc: 0.

```

Epoch 18/30
60000/60000 [=====] - 5s 89us/step - loss: 0.0700 - acc: 0.9
Epoch 19/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0695 - acc: 0.9
Epoch 20/30
60000/60000 [=====] - 5s 89us/step - loss: 0.0640 - acc: 0.9
Epoch 21/30
60000/60000 [=====] - 5s 91us/step - loss: 0.0616 - acc: 0.9
Epoch 22/30
60000/60000 [=====] - 5s 90us/step - loss: 0.0604 - acc: 0.9
Epoch 23/30
60000/60000 [=====] - 5s 89us/step - loss: 0.0577 - acc: 0.9
Epoch 24/30
60000/60000 [=====] - 5s 87us/step - loss: 0.0599 - acc: 0.9
Epoch 25/30
60000/60000 [=====] - 5s 90us/step - loss: 0.0567 - acc: 0.9
Epoch 26/30
60000/60000 [=====] - 5s 90us/step - loss: 0.0529 - acc: 0.9
Epoch 27/30
60000/60000 [=====] - 5s 89us/step - loss: 0.0540 - acc: 0.9
Epoch 28/30
60000/60000 [=====] - 5s 88us/step - loss: 0.0505 - acc: 0.9
Epoch 29/30
60000/60000 [=====] - 6s 92us/step - loss: 0.0521 - acc: 0.9
Epoch 30/30
60000/60000 [=====] - 6s 92us/step - loss: 0.0482 - acc: 0.9

```

```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_val, Y_val))

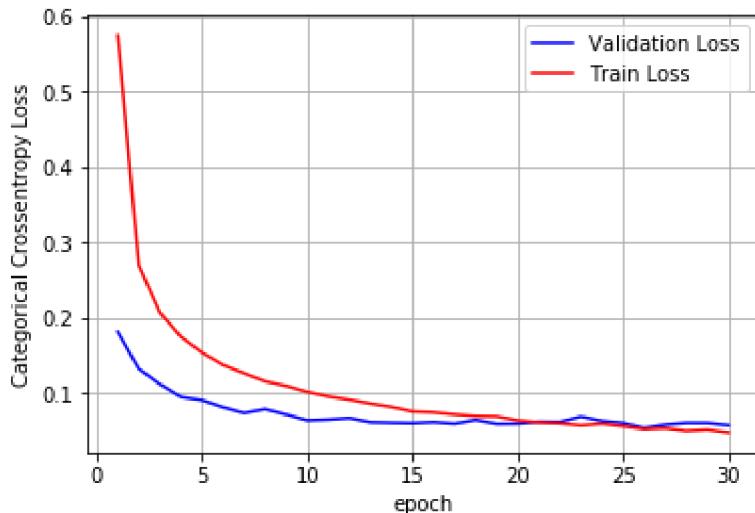
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

→ Test score: 0.05789125898362254
 Test accuracy: 0.9844



```
%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt

w_after = model.get_weights()

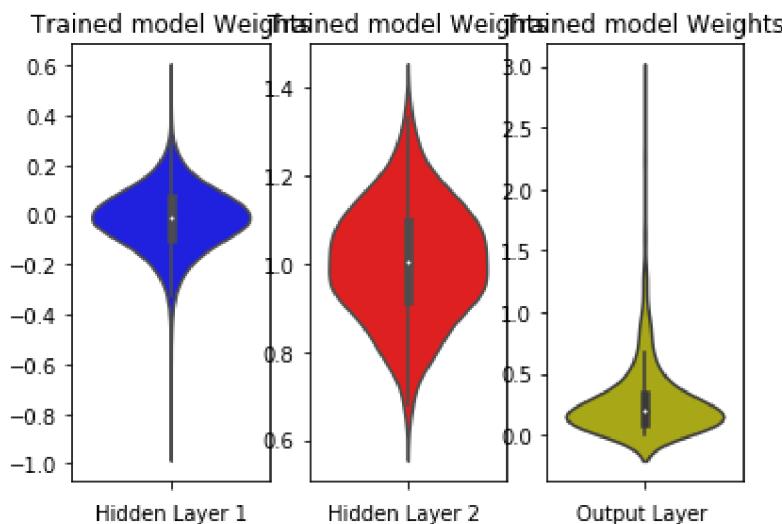
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

→



MLP + BN + Dropout + AdamOptimizer + ReLu activation

Architecture 2: 784-456-128-10;
dropout: .5,5

std = $\sqrt{2/\text{fan-in}}$ as it is ReLu

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
import warnings
warnings.filterwarnings('ignore')

model2 = Sequential()

model2.add(Dense(456, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=1))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))

model2.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=2))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))

model2.add(Dense(output_dim, activation='softmax'))
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history2 = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_
```

→

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 10s 165us/step - loss: 0.5033 - acc: 0.5

Epoch 2/30

60000/60000 [=====] - 6s 94us/step - loss: 0.2371 - acc: 0.5

Epoch 3/30

60000/60000 [=====] - 6s 94us/step - loss: 0.1830 - acc: 0.5

Epoch 4/30

60000/60000 [=====] - 6s 97us/step - loss: 0.1527 - acc: 0.5

Epoch 5/30

60000/60000 [=====] - 6s 94us/step - loss: 0.1358 - acc: 0.5

Epoch 6/30

60000/60000 [=====] - 6s 95us/step - loss: 0.1195 - acc: 0.5

Epoch 7/30

60000/60000 [=====] - 6s 96us/step - loss: 0.1078 - acc: 0.5

Epoch 8/30

60000/60000 [=====] - 6s 94us/step - loss: 0.1015 - acc: 0.5

Epoch 9/30

60000/60000 [=====] - 6s 94us/step - loss: 0.0918 - acc: 0.5

Epoch 10/30

60000/60000 [=====] - 6s 94us/step - loss: 0.0878 - acc: 0.5

Epoch 11/30

60000/60000 [=====] - 6s 93us/step - loss: 0.0861 - acc: 0.5

Epoch 12/30

60000/60000 [=====] - 6s 93us/step - loss: 0.0768 - acc: 0.5

Epoch 13/30

60000/60000 [=====] - 6s 94us/step - loss: 0.0740 - acc: 0.5

Epoch 14/30

60000/60000 [=====] - 6s 97us/step - loss: 0.0736 - acc: 0.5

Epoch 15/30

60000/60000 [=====] - 6s 96us/step - loss: 0.0678 - acc: 0.5

Epoch 16/30

60000/60000 [=====] - 6s 95us/step - loss: 0.0644 - acc: 0.5

Epoch 17/30

60000/60000 [=====] - 6s 94us/step - loss: 0.0650 - acc: 0.5

Epoch 18/30

60000/60000 [=====] - 6s 95us/step - loss: 0.0606 - acc: 0.5

Epoch 19/30

60000/60000 [=====] - 6s 93us/step - loss: 0.0605 - acc: 0.5

Epoch 20/30

60000/60000 [=====] - 6s 93us/step - loss: 0.0559 - acc: 0.5

Epoch 21/30

60000/60000 [=====] - 6s 94us/step - loss: 0.0548 - acc: 0.5

Epoch 22/30

60000/60000 [=====] - 6s 92us/step - loss: 0.0533 - acc: 0.5

Epoch 23/30

60000/60000 [=====] - 6s 96us/step - loss: 0.0511 - acc: 0.5

Epoch 24/30

60000/60000 [=====] - 6s 95us/step - loss: 0.0519 - acc: 0.5

Epoch 25/30

60000/60000 [=====] - 6s 94us/step - loss: 0.0493 - acc: 0.5

Epoch 26/30

60000/60000 [=====] - 6s 92us/step - loss: 0.0492 - acc: 0.5

Epoch 27/30

60000/60000 [=====] - 6s 95us/step - loss: 0.0459 - acc: 0.5

Epoch 28/30

60000/60000 [=====] - 6s 95us/step - loss: 0.0449 - acc: 0.5

```
Epoch 29/30
60000/60000 [=====] - 6s 98us/step - loss: 0.0430 - acc: 0.9
Epoch 30/30
60000/60000 [=====] - 6s 94us/step - loss: 0.0393 - acc: 0.9
```

```
%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

score2 = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score2[0])
print('Test accuracy:', score2[1])

fig2,ax2 = plt.subplots(1,1)
ax2.set_xlabel('epoch') ; ax2.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x2 = list(range(1,nb_epoch+1))

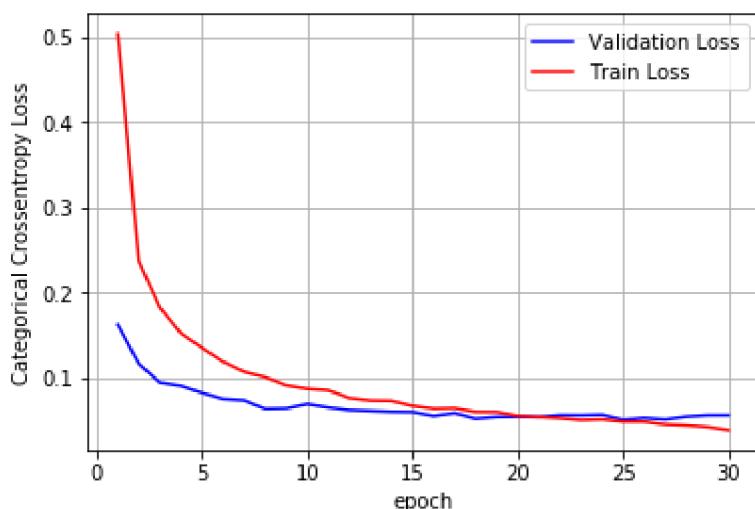
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_val, Y_val))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy2 = history2.history['val_loss']
ty2 = history2.history['loss']
plt_dynamic(x2, vy2, ty2, ax2)
```

⇒ Test score: 0.056721248715037656
 Test accuracy: 0.984



```
%matplotlib notebook
%matplotlib inline
```

```

import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

w_after2 = model2.get_weights()

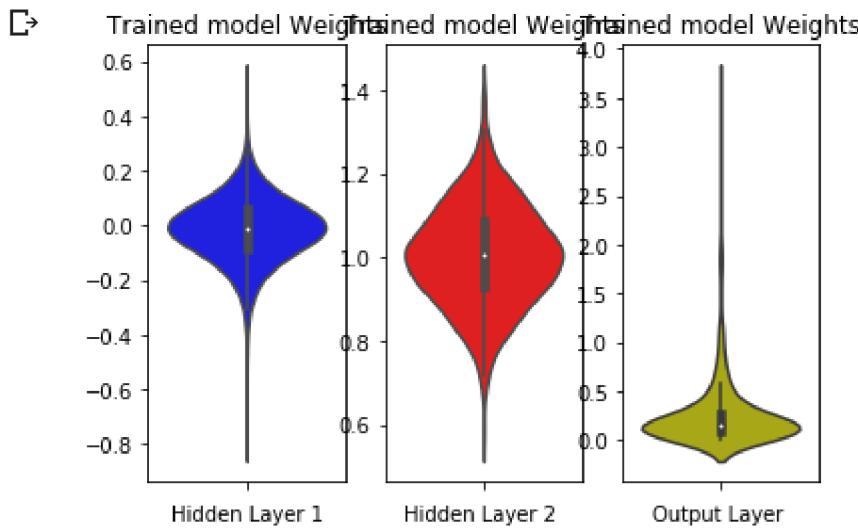
h1_w = w_after2[0].flatten().reshape(-1,1)
h2_w = w_after2[2].flatten().reshape(-1,1)
out_w = w_after2[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()

```



MLP + BN + Dropout + AdamOptimizer + ReLu activation

Architecture 3: 784-256-128-10;
dropout: .5,.5

std = $\sqrt{2/\text{fan-in}}$ as it is ReLu

```

from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
import warnings
warnings.filterwarnings('ignore')

```

```
model3 = Sequential()

model3.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=1))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=1))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(output_dim, activation='softmax'))
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history3 = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_
```

➡

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 10s 166us/step - loss: 0.6093 - acc: 0.5

Epoch 2/30

60000/60000 [=====] - 5s 89us/step - loss: 0.2919 - acc: 0.5

Epoch 3/30

60000/60000 [=====] - 5s 90us/step - loss: 0.2280 - acc: 0.5

Epoch 4/30

60000/60000 [=====] - 5s 86us/step - loss: 0.1943 - acc: 0.5

Epoch 5/30

60000/60000 [=====] - 5s 89us/step - loss: 0.1661 - acc: 0.5

Epoch 6/30

60000/60000 [=====] - 5s 89us/step - loss: 0.1532 - acc: 0.5

Epoch 7/30

60000/60000 [=====] - 5s 88us/step - loss: 0.1386 - acc: 0.5

Epoch 8/30

60000/60000 [=====] - 5s 89us/step - loss: 0.1285 - acc: 0.5

Epoch 9/30

60000/60000 [=====] - 5s 87us/step - loss: 0.1190 - acc: 0.5

Epoch 10/30

60000/60000 [=====] - 5s 88us/step - loss: 0.1156 - acc: 0.5

Epoch 11/30

60000/60000 [=====] - 5s 88us/step - loss: 0.1078 - acc: 0.5

Epoch 12/30

60000/60000 [=====] - 5s 88us/step - loss: 0.0999 - acc: 0.5

Epoch 13/30

60000/60000 [=====] - 5s 88us/step - loss: 0.0949 - acc: 0.5

Epoch 14/30

60000/60000 [=====] - 5s 87us/step - loss: 0.0927 - acc: 0.5

Epoch 15/30

60000/60000 [=====] - 5s 88us/step - loss: 0.0906 - acc: 0.5

Epoch 16/30

60000/60000 [=====] - 5s 89us/step - loss: 0.0868 - acc: 0.5

Epoch 17/30

60000/60000 [=====] - 5s 90us/step - loss: 0.0830 - acc: 0.5

Epoch 18/30

60000/60000 [=====] - 5s 87us/step - loss: 0.0791 - acc: 0.5

Epoch 19/30

60000/60000 [=====] - 5s 89us/step - loss: 0.0768 - acc: 0.5

Epoch 20/30

60000/60000 [=====] - 5s 88us/step - loss: 0.0761 - acc: 0.5

Epoch 21/30

60000/60000 [=====] - 5s 89us/step - loss: 0.0716 - acc: 0.5

Epoch 22/30

60000/60000 [=====] - 5s 86us/step - loss: 0.0721 - acc: 0.5

Epoch 23/30

60000/60000 [=====] - 5s 87us/step - loss: 0.0696 - acc: 0.5

Epoch 24/30

60000/60000 [=====] - 5s 87us/step - loss: 0.0663 - acc: 0.5

Epoch 25/30

60000/60000 [=====] - 5s 89us/step - loss: 0.0662 - acc: 0.5

Epoch 26/30

60000/60000 [=====] - 5s 88us/step - loss: 0.0630 - acc: 0.5

Epoch 27/30

60000/60000 [=====] - 5s 88us/step - loss: 0.0594 - acc: 0.5

Epoch 28/30

60000/60000 [=====] - 5s 88us/step - loss: 0.0605 - acc: 0.5

```

Epoch 29/30
60000/60000 [=====] - 5s 90us/step - loss: 0.0582 - acc: 0.9
Epoch 30/30
60000/60000 [=====] - 5s 87us/step - loss: 0.0564 - acc: 0.9

```

```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

score3 = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score3[0])
print('Test accuracy:', score3[1])

fig3,ax3 = plt.subplots(1,1)
ax3.set_xlabel('epoch') ; ax3.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x3 = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_val, Y_val))

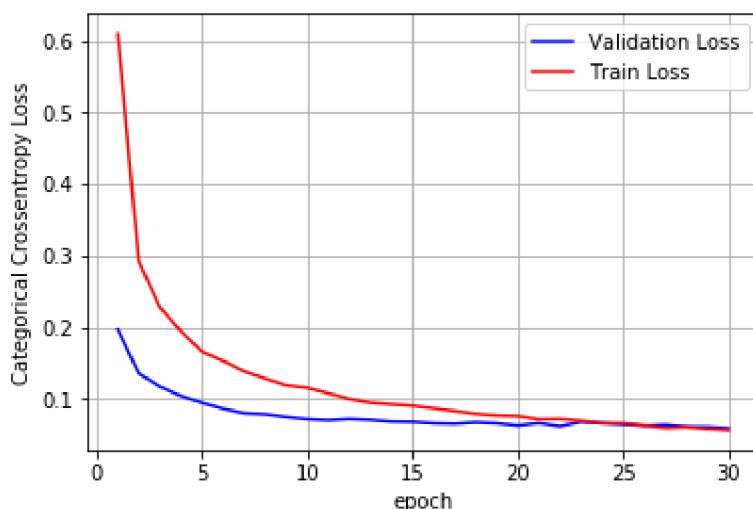
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy3 = history3.history['val_loss']
ty3 = history3.history['loss']
plt_dynamic(x3, vy3, ty3, ax3)

```

⇒ Test score: 0.05895621230599354
 Test accuracy: 0.9823



MLP + BN + Dropout + AdamOptimizer + ReLu activation

Architecture 4: 784-256-96-10;
dropout: .5,.5

std = $\sqrt{2/\text{fan-in}}$ as it is ReLu

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
import warnings
warnings.filterwarnings('ignore')

model4 = Sequential()

model4.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.14433, seed=1))
model4.add(BatchNormalization())
model4.add(Dropout(0.5))

model4.add(Dense(96, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.14433, seed=2))
model4.add(BatchNormalization())
model4.add(Dropout(0.5))

model4.add(Dense(output_dim, activation='softmax'))
model4.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history4 = model4.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_
```

⇨

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 10s 168us/step - loss: 0.6637 - acc: 0.5

Epoch 2/30

60000/60000 [=====] - 5s 91us/step - loss: 0.3106 - acc: 0.5

Epoch 3/30

60000/60000 [=====] - 5s 87us/step - loss: 0.2386 - acc: 0.5

Epoch 4/30

60000/60000 [=====] - 5s 88us/step - loss: 0.2054 - acc: 0.5

Epoch 5/30

60000/60000 [=====] - 5s 91us/step - loss: 0.1773 - acc: 0.5

Epoch 6/30

60000/60000 [=====] - 5s 91us/step - loss: 0.1587 - acc: 0.5

Epoch 7/30

60000/60000 [=====] - 5s 90us/step - loss: 0.1454 - acc: 0.5

Epoch 8/30

60000/60000 [=====] - 5s 90us/step - loss: 0.1385 - acc: 0.5

Epoch 9/30

60000/60000 [=====] - 5s 88us/step - loss: 0.1274 - acc: 0.5

Epoch 10/30

60000/60000 [=====] - 5s 89us/step - loss: 0.1181 - acc: 0.5

Epoch 11/30

60000/60000 [=====] - 5s 89us/step - loss: 0.1108 - acc: 0.5

Epoch 12/30

60000/60000 [=====] - 5s 89us/step - loss: 0.1060 - acc: 0.5

Epoch 13/30

60000/60000 [=====] - 5s 90us/step - loss: 0.1036 - acc: 0.5

Epoch 14/30

60000/60000 [=====] - 5s 89us/step - loss: 0.0976 - acc: 0.5

Epoch 15/30

60000/60000 [=====] - 5s 90us/step - loss: 0.0956 - acc: 0.5

Epoch 16/30

60000/60000 [=====] - 5s 89us/step - loss: 0.0915 - acc: 0.5

Epoch 17/30

60000/60000 [=====] - 6s 93us/step - loss: 0.0862 - acc: 0.5

Epoch 18/30

60000/60000 [=====] - 5s 89us/step - loss: 0.0829 - acc: 0.5

Epoch 19/30

60000/60000 [=====] - 5s 89us/step - loss: 0.0803 - acc: 0.5

Epoch 20/30

60000/60000 [=====] - 5s 92us/step - loss: 0.0772 - acc: 0.5

Epoch 21/30

60000/60000 [=====] - 5s 88us/step - loss: 0.0763 - acc: 0.5

Epoch 22/30

60000/60000 [=====] - 5s 91us/step - loss: 0.0753 - acc: 0.5

Epoch 23/30

60000/60000 [=====] - 6s 93us/step - loss: 0.0710 - acc: 0.5

Epoch 24/30

60000/60000 [=====] - 5s 90us/step - loss: 0.0710 - acc: 0.5

Epoch 25/30

60000/60000 [=====] - 5s 87us/step - loss: 0.0730 - acc: 0.5

Epoch 26/30

60000/60000 [=====] - 5s 89us/step - loss: 0.0659 - acc: 0.5

Epoch 27/30

60000/60000 [=====] - 5s 90us/step - loss: 0.0634 - acc: 0.5

Epoch 28/30

60000/60000 [=====] - 5s 88us/step - loss: 0.0674 - acc: 0.5

```

Epoch 29/30
60000/60000 [=====] - 5s 87us/step - loss: 0.0611 - acc: 0.9
Epoch 30/30
60000/60000 [=====] - 5s 90us/step - loss: 0.0598 - acc: 0.9

```

```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

score4 = model4.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score4[0])
print('Test accuracy:', score4[1])

fig4,ax4 = plt.subplots(1,1)
ax4.set_xlabel('epoch') ; ax4.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x4 = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_val, Y_val))

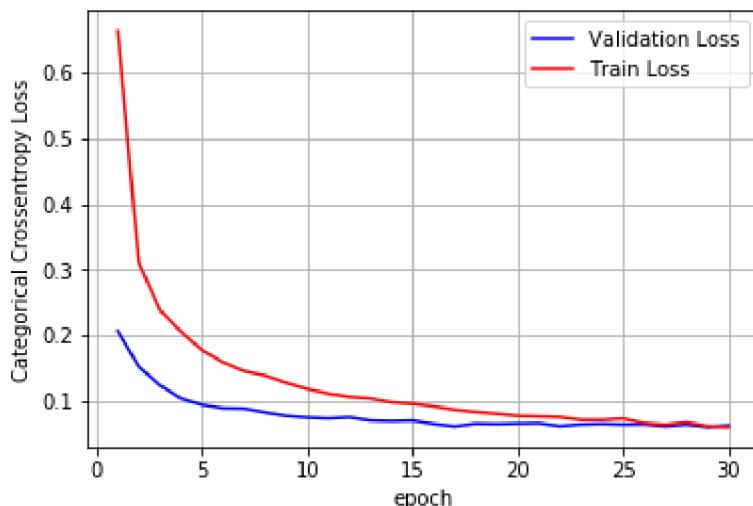
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy4 = history4.history['val_loss']
ty4 = history4.history['loss']
plt_dynamic(x4, vy4, ty4, ax4)

```

⇒ Test score: 0.06231452802268032
 Test accuracy: 0.9832



MLP + BN + Dropout + AdamOptimizer + ReLu activation

Part 2 : Here we use 3 Hidden Layers

Architecture 5: 784-512-392-128-10;
dropout: .5,.5

std = $\sqrt{2/\text{fan-in}}$ as it is ReLu

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
import warnings
warnings.filterwarnings('ignore')

model5 = Sequential()

model5.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.07142, seed=1))
model5.add(BatchNormalization())
model5.add(Dropout(0.5))

model5.add(Dense(392, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=2))
model5.add(BatchNormalization())
model5.add(Dropout(0.5))

model5.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=3))
model5.add(BatchNormalization())
model5.add(Dropout(0.5))

model5.add(Dense(output_dim, activation='softmax'))
model5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model5.summary()

history5 = model5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_
```



Model: "sequential_18"

Layer (type)	Output Shape	Param #
dense_56 (Dense)	(None, 512)	401920
batch_normalization_39 (Batch Normalization)	(None, 512)	2048
dropout_39 (Dropout)	(None, 512)	0
dense_57 (Dense)	(None, 392)	201096
batch_normalization_40 (Batch Normalization)	(None, 392)	1568
dropout_40 (Dropout)	(None, 392)	0
dense_58 (Dense)	(None, 128)	50304
batch_normalization_41 (Batch Normalization)	(None, 128)	512
dropout_41 (Dropout)	(None, 128)	0
dense_59 (Dense)	(None, 10)	1290

Total params: 658,738

Trainable params: 656,674

Non-trainable params: 2,064

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 13s 212us/step - loss: 0.5931 - acc: 0.

Epoch 2/30

60000/60000 [=====] - 7s 123us/step - loss: 0.2522 - acc: 0.

Epoch 3/30

60000/60000 [=====] - 7s 121us/step - loss: 0.1926 - acc: 0.

Epoch 4/30

60000/60000 [=====] - 7s 123us/step - loss: 0.1629 - acc: 0.

Epoch 5/30

60000/60000 [=====] - 7s 123us/step - loss: 0.1417 - acc: 0.

Epoch 6/30

60000/60000 [=====] - 7s 122us/step - loss: 0.1293 - acc: 0.

Epoch 7/30

60000/60000 [=====] - 7s 122us/step - loss: 0.1146 - acc: 0.

Epoch 8/30

60000/60000 [=====] - 7s 122us/step - loss: 0.1092 - acc: 0.

Epoch 9/30

60000/60000 [=====] - 7s 121us/step - loss: 0.1003 - acc: 0.

Epoch 10/30

60000/60000 [=====] - 7s 123us/step - loss: 0.0945 - acc: 0.

Epoch 11/30

60000/60000 [=====] - 7s 120us/step - loss: 0.0888 - acc: 0.

Epoch 12/30

60000/60000 [=====] - 7s 120us/step - loss: 0.0827 - acc: 0.

Epoch 13/30

60000/60000 [=====] - 7s 121us/step - loss: 0.0810 - acc: 0.

Epoch 14/30

60000/60000 [=====] - 7s 121us/step - loss: 0.0770 - acc: 0.

```

Epoch 15/30
60000/60000 [=====] - 7s 121us/step - loss: 0.0727 - acc: 0.
Epoch 16/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0682 - acc: 0.
Epoch 17/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0703 - acc: 0.
Epoch 18/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0641 - acc: 0.
Epoch 19/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0604 - acc: 0.
Epoch 20/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0624 - acc: 0.
Epoch 21/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0586 - acc: 0.
Epoch 22/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0543 - acc: 0.
Epoch 23/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0510 - acc: 0.
Epoch 24/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0549 - acc: 0.
Epoch 25/30
60000/60000 [=====] - 7s 124us/step - loss: 0.0511 - acc: 0.
Epoch 26/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0496 - acc: 0.
Epoch 27/30
60000/60000 [=====] - 7s 121us/step - loss: 0.0495 - acc: 0.
Epoch 28/30
60000/60000 [=====] - 7s 121us/step - loss: 0.0498 - acc: 0.
Epoch 29/30
60000/60000 [=====] - 7s 121us/step - loss: 0.0449 - acc: 0.
Epoch 30/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0442 - acc: 0.

```

```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

score5 = model5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score5[0])
print('Test accuracy:', score5[1])

fig5,ax5 = plt.subplots(1,1)
ax5.set_xlabel('epoch') ; ax5.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x5 = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data
# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

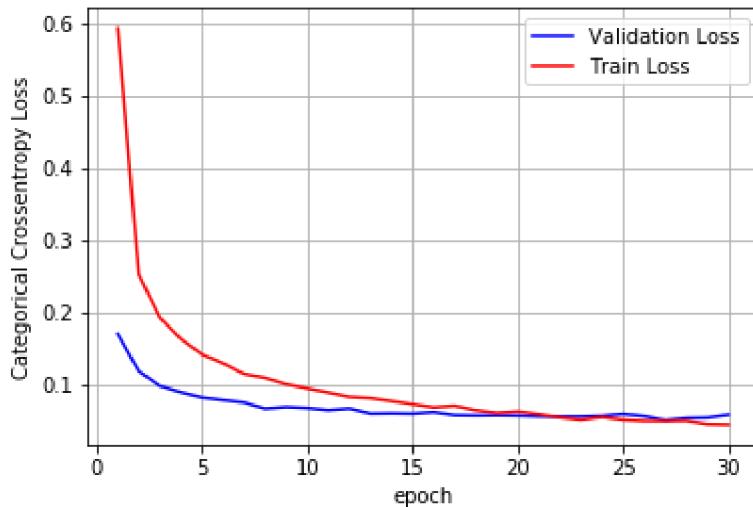
```

```
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
```

```
vy5 = history5.history['val_loss']
ty5 = history5.history['loss']
plt_dynamic(x5, vy5, ty5, ax5)
```

→ Test score: 0.057912785604948294

Test accuracy: 0.9846



```
%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

w_after5 = model5.get_weights()

h1_w = w_after5[0].flatten().reshape(-1,1)
h2_w = w_after5[2].flatten().reshape(-1,1)
h3_w = w_after5[4].flatten().reshape(-1,1)
out_w = w_after5[6].flatten().reshape(-1,1)

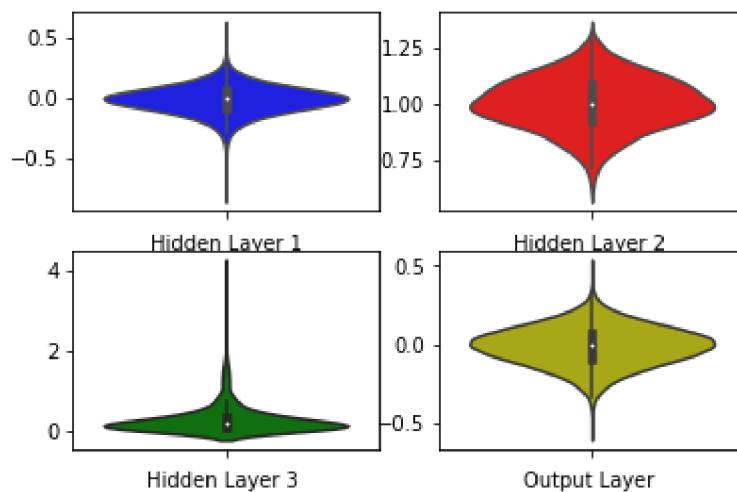
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(2, 2, 1)
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(2, 2, 2)
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(2, 2, 3)
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(2, 2, 4)
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()
```

→



MLP + BN + Dropout + AdamOptimizer + ReLu activation

Part 2 : Here we use 3 Hidden Layers

Architecture 6: 784-600-300-100-10;
dropout: .5,.5

`std = sqrt(2/fan-in)` as it is ReLu

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
import warnings
warnings.filterwarnings('ignore')

model6 = Sequential()

model6.add(Dense(600, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.08164, see
model6.add(BatchNormalization())
model6.add(Dropout(0.5))

model6.add(Dense(300, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.14142, see
model6.add(BatchNormalization())
model6.add(Dropout(0.5))

model6.add(Dense(100, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.28284, see
model6.add(BatchNormalization())
model6.add(Dropout(0.5))

model6.add(Dense(output_dim, activation='softmax'))
model6.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model6.summary()
```

```
history6 = model6.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_
```



Model: "sequential_19"

Layer (type)	Output Shape	Param #
dense_60 (Dense)	(None, 600)	471000
batch_normalization_42 (Batch Normalization)	(None, 600)	2400
dropout_42 (Dropout)	(None, 600)	0
dense_61 (Dense)	(None, 300)	180300
batch_normalization_43 (Batch Normalization)	(None, 300)	1200
dropout_43 (Dropout)	(None, 300)	0
dense_62 (Dense)	(None, 100)	30100
batch_normalization_44 (Batch Normalization)	(None, 100)	400
dropout_44 (Dropout)	(None, 100)	0
dense_63 (Dense)	(None, 10)	1010

Total params: 686,410
 Trainable params: 684,410
 Non-trainable params: 2,000

Train on 60000 samples, validate on 10000 samples

Epoch 1/30
 60000/60000 [=====] - 13s 221us/step - loss: 0.6130 - acc: 0.
 Epoch 2/30
 60000/60000 [=====] - 8s 125us/step - loss: 0.2548 - acc: 0.
 Epoch 3/30
 60000/60000 [=====] - 8s 126us/step - loss: 0.1962 - acc: 0.
 Epoch 4/30
 60000/60000 [=====] - 7s 124us/step - loss: 0.1651 - acc: 0.
 Epoch 5/30
 60000/60000 [=====] - 8s 131us/step - loss: 0.1409 - acc: 0.
 Epoch 6/30
 60000/60000 [=====] - 8s 126us/step - loss: 0.1297 - acc: 0.
 Epoch 7/30
 60000/60000 [=====] - 7s 124us/step - loss: 0.1151 - acc: 0.
 Epoch 8/30
 60000/60000 [=====] - 8s 126us/step - loss: 0.1085 - acc: 0.
 Epoch 9/30
 60000/60000 [=====] - 8s 125us/step - loss: 0.1015 - acc: 0.
 Epoch 10/30
 60000/60000 [=====] - 8s 126us/step - loss: 0.0937 - acc: 0.
 Epoch 11/30
 60000/60000 [=====] - 8s 127us/step - loss: 0.0904 - acc: 0.
 Epoch 12/30
 60000/60000 [=====] - 8s 125us/step - loss: 0.0829 - acc: 0.
 Epoch 13/30
 60000/60000 [=====] - 7s 125us/step - loss: 0.0788 - acc: 0.
 Epoch 14/30
 60000/60000 [=====] - 8s 125us/step - loss: 0.0786 - acc: 0.

```

Epoch 15/30
60000/60000 [=====] - 7s 125us/step - loss: 0.0726 - acc: 0.
Epoch 16/30
60000/60000 [=====] - 7s 124us/step - loss: 0.0678 - acc: 0.
Epoch 17/30
60000/60000 [=====] - 8s 125us/step - loss: 0.0680 - acc: 0.
Epoch 18/30
60000/60000 [=====] - 8s 126us/step - loss: 0.0638 - acc: 0.
Epoch 19/30
60000/60000 [=====] - 7s 125us/step - loss: 0.0624 - acc: 0.
Epoch 20/30
60000/60000 [=====] - 8s 125us/step - loss: 0.0603 - acc: 0.
Epoch 21/30
60000/60000 [=====] - 8s 126us/step - loss: 0.0574 - acc: 0.
Epoch 22/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0565 - acc: 0.
Epoch 23/30
60000/60000 [=====] - 7s 125us/step - loss: 0.0537 - acc: 0.
Epoch 24/30
60000/60000 [=====] - 7s 125us/step - loss: 0.0537 - acc: 0.
Epoch 25/30
60000/60000 [=====] - 7s 125us/step - loss: 0.0503 - acc: 0.
Epoch 26/30
60000/60000 [=====] - 7s 124us/step - loss: 0.0482 - acc: 0.
Epoch 27/30
60000/60000 [=====] - 7s 124us/step - loss: 0.0491 - acc: 0.
Epoch 28/30
60000/60000 [=====] - 7s 124us/step - loss: 0.0480 - acc: 0.
Epoch 29/30
60000/60000 [=====] - 8s 127us/step - loss: 0.0461 - acc: 0.
Epoch 30/30
60000/60000 [=====] - 8s 126us/step - loss: 0.0437 - acc: 0.

```

```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

score6 = model6.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score6[0])
print('Test accuracy:', score6[1])

fig6,ax6 = plt.subplots(1,1)
ax6.set_xlabel('epoch') ; ax6.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x6 = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_val, Y_val))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

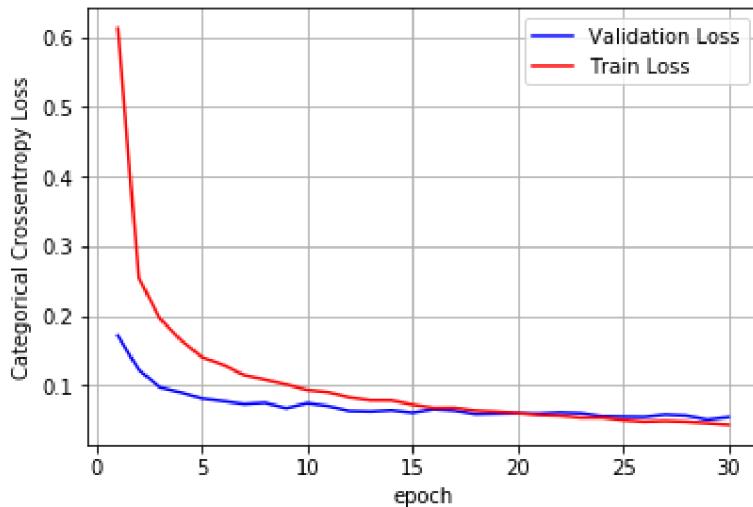
```

```
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy6 = history6.history['val_loss']
ty6 = history6.history['loss']
plt_dynamic(x6, vy6, ty6, ax6.)
```

↳ Test score: 0.05445547826328548

Test accuracy: 0.9857



MLP + BN + Dropout + AdamOptimizer + ReLu activation

Part 2 : Here we use 3 Hidden Layers

Architecture 7: 784-512-456-128-10;

dropout: .5,.5

std = $\sqrt{2/\text{fan-in}}$ as it is ReLu

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
import warnings
warnings.filterwarnings('ignore')

model7 = Sequential()

model7.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.0662, seed=1))
model7.add(BatchNormalization())
model7.add(Dropout(0.5))

model7.add(Dense(456, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.0662, seed=2))
model7.add(BatchNormalization())
model7.add(Dropout(0.5))

model7.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=3))
model7.add(BatchNormalization())
model7.add(Dropout(0.5))

model7.add(Dense(output_dim, activation='softmax'))
model7.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
model7.summary()
```

```
history7 = model7.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_
```

⇨

Model: "sequential_20"

Layer (type)	Output Shape	Param #
dense_64 (Dense)	(None, 512)	401920
batch_normalization_45 (Batch Normalization)	(None, 512)	2048
dropout_45 (Dropout)	(None, 512)	0
dense_65 (Dense)	(None, 456)	233928
batch_normalization_46 (Batch Normalization)	(None, 456)	1824
dropout_46 (Dropout)	(None, 456)	0
dense_66 (Dense)	(None, 128)	58496
batch_normalization_47 (Batch Normalization)	(None, 128)	512
dropout_47 (Dropout)	(None, 128)	0
dense_67 (Dense)	(None, 10)	1290

Total params: 700,018
 Trainable params: 697,826
 Non-trainable params: 2,192

Train on 60000 samples, validate on 10000 samples

Epoch 1/30
 60000/60000 [=====] - 14s 227us/step - loss: 0.5690 - acc: 0.
 Epoch 2/30
 60000/60000 [=====] - 7s 123us/step - loss: 0.2451 - acc: 0.
 Epoch 3/30
 60000/60000 [=====] - 7s 124us/step - loss: 0.1847 - acc: 0.
 Epoch 4/30
 60000/60000 [=====] - 7s 122us/step - loss: 0.1569 - acc: 0.
 Epoch 5/30
 60000/60000 [=====] - 7s 124us/step - loss: 0.1392 - acc: 0.
 Epoch 6/30
 60000/60000 [=====] - 7s 122us/step - loss: 0.1233 - acc: 0.
 Epoch 7/30
 60000/60000 [=====] - 7s 122us/step - loss: 0.1084 - acc: 0.
 Epoch 8/30
 60000/60000 [=====] - 7s 123us/step - loss: 0.1011 - acc: 0.
 Epoch 9/30
 60000/60000 [=====] - 7s 123us/step - loss: 0.0986 - acc: 0.
 Epoch 10/30
 60000/60000 [=====] - 7s 122us/step - loss: 0.0933 - acc: 0.
 Epoch 11/30
 60000/60000 [=====] - 7s 122us/step - loss: 0.0850 - acc: 0.
 Epoch 12/30
 60000/60000 [=====] - 7s 121us/step - loss: 0.0817 - acc: 0.
 Epoch 13/30
 60000/60000 [=====] - 7s 121us/step - loss: 0.0780 - acc: 0.
 Epoch 14/30
 60000/60000 [=====] - 7s 123us/step - loss: 0.0756 - acc: 0.

```

Epoch 15/30
60000/60000 [=====] - 7s 125us/step - loss: 0.0727 - acc: 0.
Epoch 16/30
60000/60000 [=====] - 8s 127us/step - loss: 0.0677 - acc: 0.
Epoch 17/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0653 - acc: 0.
Epoch 18/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0615 - acc: 0.
Epoch 19/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0615 - acc: 0.
Epoch 20/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0593 - acc: 0.
Epoch 21/30
60000/60000 [=====] - 7s 124us/step - loss: 0.0517 - acc: 0.
Epoch 22/30
60000/60000 [=====] - 7s 121us/step - loss: 0.0572 - acc: 0.
Epoch 23/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0511 - acc: 0.
Epoch 24/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0517 - acc: 0.
Epoch 25/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0496 - acc: 0.
Epoch 26/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0474 - acc: 0.
Epoch 27/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0478 - acc: 0.
Epoch 28/30
60000/60000 [=====] - 7s 123us/step - loss: 0.0457 - acc: 0.
Epoch 29/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0411 - acc: 0.
Epoch 30/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0454 - acc: 0.

```

```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

score7 = model7.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score7[0])
print('Test accuracy:', score7[1])

fig7,ax7 = plt.subplots(1,1)
ax7.set_xlabel('epoch') ; ax7.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x7 = list(range(1,nb_epoch+1))

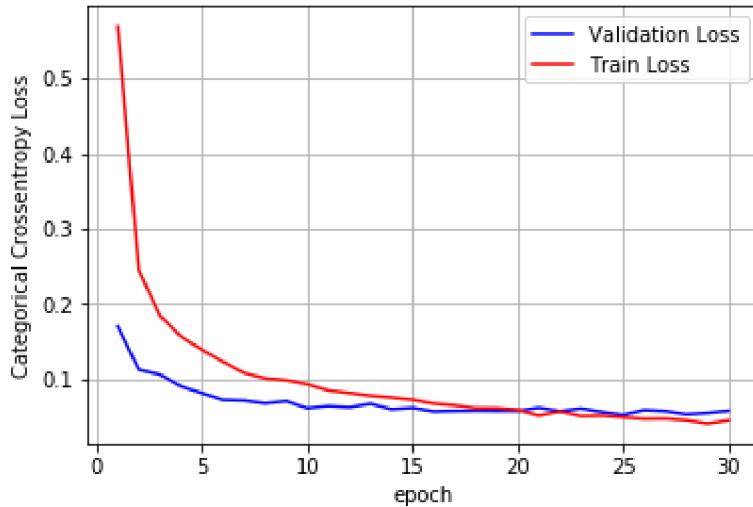
vy7 = history7.history['val_loss']
ty7 = history7.history['loss']
plt_dynamic(x7, vy7, ty7, ax7)

```



Test score: 0.05791448878053052

Test accuracy: 0.9839



MLP + BN + Dropout + AdamOptimizer + ReLu activation

Part 2 : Here we use 3 Hidden Layers

Architecture 8: 784-456-128-32-10;
dropout: .5,.5

std = $\sqrt{2/\text{fan-in}}$ as it is ReLu

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
import warnings
warnings.filterwarnings('ignore')

model8 = Sequential()

model8.add(Dense(456, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.125, seed=No
model8.add(BatchNormalization())
model8.add(Dropout(0.5))

model8.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.125, seed=No
model8.add(BatchNormalization())
model8.add(Dropout(0.5))

model8.add(Dense(32, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.25, seed=No
model8.add(BatchNormalization())
model8.add(Dropout(0.5))

model8.add(Dense(output_dim, activation='softmax'))
model8.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model8.summary()

history8 = model8.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_
```



Model: "sequential_21"

Layer (type)	Output Shape	Param #
dense_68 (Dense)	(None, 456)	357960
batch_normalization_48 (Batch Normalization)	(None, 456)	1824
dropout_48 (Dropout)	(None, 456)	0
dense_69 (Dense)	(None, 128)	58496
batch_normalization_49 (Batch Normalization)	(None, 128)	512
dropout_49 (Dropout)	(None, 128)	0
dense_70 (Dense)	(None, 32)	4128
batch_normalization_50 (Batch Normalization)	(None, 32)	128
dropout_50 (Dropout)	(None, 32)	0
dense_71 (Dense)	(None, 10)	330

Total params: 423,378
 Trainable params: 422,146
 Non-trainable params: 1,232

Train on 60000 samples, validate on 10000 samples

Epoch 1/30
 60000/60000 [=====] - 14s 232us/step - loss: 0.9679 - acc: 0.
 Epoch 2/30
 60000/60000 [=====] - 7s 120us/step - loss: 0.4291 - acc: 0.
 Epoch 3/30
 60000/60000 [=====] - 7s 121us/step - loss: 0.3162 - acc: 0.
 Epoch 4/30
 60000/60000 [=====] - 7s 120us/step - loss: 0.2564 - acc: 0.
 Epoch 5/30
 60000/60000 [=====] - 7s 118us/step - loss: 0.2209 - acc: 0.
 Epoch 6/30
 60000/60000 [=====] - 7s 119us/step - loss: 0.1987 - acc: 0.
 Epoch 7/30
 60000/60000 [=====] - 7s 117us/step - loss: 0.1805 - acc: 0.
 Epoch 8/30
 60000/60000 [=====] - 7s 121us/step - loss: 0.1696 - acc: 0.
 Epoch 9/30
 60000/60000 [=====] - 7s 118us/step - loss: 0.1608 - acc: 0.
 Epoch 10/30
 60000/60000 [=====] - 7s 119us/step - loss: 0.1469 - acc: 0.
 Epoch 11/30
 60000/60000 [=====] - 7s 121us/step - loss: 0.1432 - acc: 0.
 Epoch 12/30
 60000/60000 [=====] - 7s 118us/step - loss: 0.1295 - acc: 0.
 Epoch 13/30
 60000/60000 [=====] - 7s 123us/step - loss: 0.1219 - acc: 0.
 Epoch 14/30
 60000/60000 [=====] - 7s 120us/step - loss: 0.1233 - acc: 0.

```

Epoch 15/30
60000/60000 [=====] - 7s 119us/step - loss: 0.1133 - acc: 0.
Epoch 16/30
60000/60000 [=====] - 7s 118us/step - loss: 0.1098 - acc: 0.
Epoch 17/30
60000/60000 [=====] - 7s 120us/step - loss: 0.1049 - acc: 0.
Epoch 18/30
60000/60000 [=====] - 7s 119us/step - loss: 0.1025 - acc: 0.
Epoch 19/30
60000/60000 [=====] - 7s 121us/step - loss: 0.0978 - acc: 0.
Epoch 20/30
60000/60000 [=====] - 7s 120us/step - loss: 0.0947 - acc: 0.
Epoch 21/30
60000/60000 [=====] - 7s 120us/step - loss: 0.0945 - acc: 0.
Epoch 22/30
60000/60000 [=====] - 7s 120us/step - loss: 0.0883 - acc: 0.
Epoch 23/30
60000/60000 [=====] - 7s 117us/step - loss: 0.0850 - acc: 0.
Epoch 24/30
60000/60000 [=====] - 7s 119us/step - loss: 0.0873 - acc: 0.
Epoch 25/30
60000/60000 [=====] - 7s 120us/step - loss: 0.0792 - acc: 0.
Epoch 26/30
60000/60000 [=====] - 7s 119us/step - loss: 0.0801 - acc: 0.
Epoch 27/30
60000/60000 [=====] - 7s 124us/step - loss: 0.0806 - acc: 0.
Epoch 28/30
60000/60000 [=====] - 7s 122us/step - loss: 0.0779 - acc: 0.
Epoch 29/30
60000/60000 [=====] - 7s 120us/step - loss: 0.0750 - acc: 0.
Epoch 30/30
60000/60000 [=====] - 7s 121us/step - loss: 0.0759 - acc: 0.

```

```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

score8 = model8.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score8[0])
print('Test accuracy:', score8[1])

fig8,ax8 = plt.subplots(1,1)
ax8.set_xlabel('epoch') ; ax8.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x8 = list(range(1,nb_epoch+1))

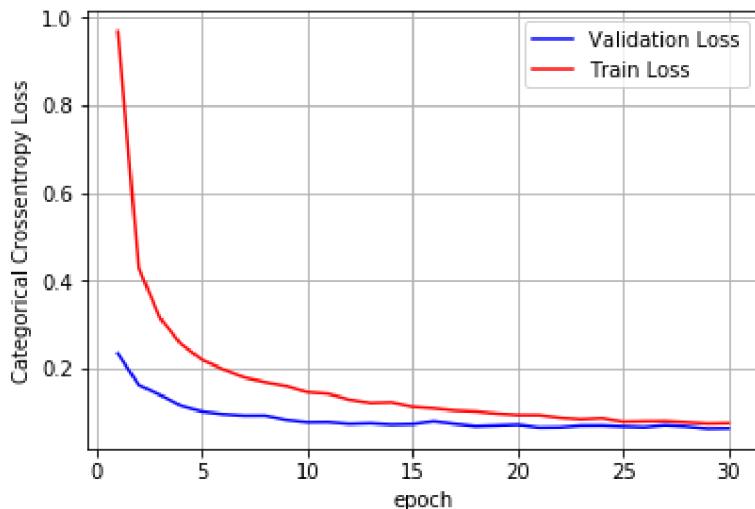
vy8 = history8.history['val_loss']
ty8 = history8.history['loss']
plt_dynamic(x8, vy8, ty8, ax8).

```



Test score: 0.06451866653093893

Test accuracy: 0.9845



MLP + BN + Dropout + AdamOptimizer + ReLu activation

Part 3 : Here we use 5 Hidden Layers

Architecture 9: 784-600-500-400-200-128-10;
dropout: .5,.5

std = $\text{sqrt}(2/\text{fan-in})$ as it is ReLu

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
import warnings
warnings.filterwarnings('ignore')

model9 = Sequential()

model9.add(Dense(600, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.06324, seed=No))
model9.add(BatchNormalization())
model9.add(Dropout(0.5))

model9.add(Dense(500, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.0707, seed=No))
model9.add(BatchNormalization())
model9.add(Dropout(0.5))

model9.add(Dense(400, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.08164, seed=No))
model9.add(BatchNormalization())
model9.add(Dropout(0.5))

model9.add(Dense(300, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=No))
model9.add(BatchNormalization())
model9.add(Dropout(0.5))

model9.add(Dense(200, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=No))
model9.add(BatchNormalization())
model9.add(Dropout(0.5))

model9.add(Dense(output_dim, activation='softmax'))
model9.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
model9.summary()
```

```
history9 = model9.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_
```

⇨

Model: "sequential_22"

Layer (type)	Output Shape	Param #
dense_72 (Dense)	(None, 600)	471000
batch_normalization_51 (Batch Normalization)	(None, 600)	2400
dropout_51 (Dropout)	(None, 600)	0
dense_73 (Dense)	(None, 500)	300500
batch_normalization_52 (Batch Normalization)	(None, 500)	2000
dropout_52 (Dropout)	(None, 500)	0
dense_74 (Dense)	(None, 400)	200400
batch_normalization_53 (Batch Normalization)	(None, 400)	1600
dropout_53 (Dropout)	(None, 400)	0
dense_75 (Dense)	(None, 300)	120300
batch_normalization_54 (Batch Normalization)	(None, 300)	1200
dropout_54 (Dropout)	(None, 300)	0
dense_76 (Dense)	(None, 200)	60200
batch_normalization_55 (Batch Normalization)	(None, 200)	800
dropout_55 (Dropout)	(None, 200)	0
dense_77 (Dense)	(None, 10)	2010

Total params: 1,162,410

Trainable params: 1,158,410

Non-trainable params: 4,000

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 19s 312us/step - loss: 0.9268 - acc: 0

Epoch 2/30

60000/60000 [=====] - 11s 183us/step - loss: 0.3195 - acc: 0

Epoch 3/30

60000/60000 [=====] - 11s 182us/step - loss: 0.2385 - acc: 0

Epoch 4/30

60000/60000 [=====] - 11s 183us/step - loss: 0.1976 - acc: 0

Epoch 5/30

60000/60000 [=====] - 11s 184us/step - loss: 0.1706 - acc: 0

Epoch 6/30

60000/60000 [=====] - 11s 183us/step - loss: 0.1541 - acc: 0

Epoch 7/30

60000/60000 [=====] - 11s 184us/step - loss: 0.1441 - acc: 0

Epoch 8/30

60000/60000 [=====] - 11s 184us/step - loss: 0.1272 - acc: 0

```

Epoch 9/30
60000/60000 [=====] - 11s 182us/step - loss: 0.1234 - acc: 0
Epoch 10/30
60000/60000 [=====] - 11s 183us/step - loss: 0.1179 - acc: 0
Epoch 11/30
60000/60000 [=====] - 11s 183us/step - loss: 0.1090 - acc: 0
Epoch 12/30
60000/60000 [=====] - 11s 182us/step - loss: 0.1054 - acc: 0
Epoch 13/30
60000/60000 [=====] - 11s 182us/step - loss: 0.0987 - acc: 0
Epoch 14/30
60000/60000 [=====] - 11s 183us/step - loss: 0.0941 - acc: 0
Epoch 15/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0886 - acc: 0
Epoch 16/30
60000/60000 [=====] - 11s 185us/step - loss: 0.0853 - acc: 0
Epoch 17/30
60000/60000 [=====] - 11s 183us/step - loss: 0.0838 - acc: 0
Epoch 18/30
60000/60000 [=====] - 11s 182us/step - loss: 0.0780 - acc: 0
Epoch 19/30
60000/60000 [=====] - 11s 180us/step - loss: 0.0775 - acc: 0
Epoch 20/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0749 - acc: 0
Epoch 21/30
60000/60000 [=====] - 11s 185us/step - loss: 0.0706 - acc: 0
Epoch 22/30
60000/60000 [=====] - 11s 182us/step - loss: 0.0685 - acc: 0
Epoch 23/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0655 - acc: 0
Epoch 24/30
60000/60000 [=====] - 11s 183us/step - loss: 0.0651 - acc: 0
Epoch 25/30
60000/60000 [=====] - 11s 186us/step - loss: 0.0622 - acc: 0
Epoch 26/30
60000/60000 [=====] - 11s 188us/step - loss: 0.0589 - acc: 0
Epoch 27/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0581 - acc: 0
Epoch 28/30
60000/60000 [=====] - 11s 182us/step - loss: 0.0583 - acc: 0
Epoch 29/30
60000/60000 [=====] - 11s 182us/step - loss: 0.0565 - acc: 0
Epoch 30/30
60000/60000 [=====] - 11s 184us/step - loss: 0.0537 - acc: 0

```

```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

score9 = model9.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score9[0])
print('Test accuracy:', score9[1])

```

```

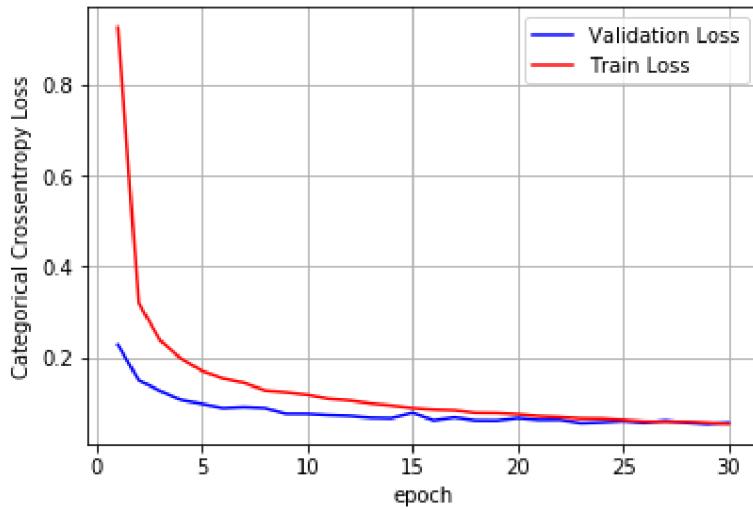
fig9,ax9 = plt.subplots(1,1)
ax9.set_xlabel('epoch') ; ax9.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x9 = list(range(1,nb_epoch+1))

vy9 = history9.history['val_loss']
ty9 = history9.history['loss']
plt_dynamic(x9, vy9, ty9, ax9)

```

↳ Test score: 0.05667422412817832
 Test accuracy: 0.9854



```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

w_after9 = model9.get_weights()

h1_w = w_after9[0].flatten().reshape(-1,1)
h2_w = w_after9[2].flatten().reshape(-1,1)
h3_w = w_after9[4].flatten().reshape(-1,1)
h4_w = w_after9[6].flatten().reshape(-1,1)
h5_w = w_after9[8].flatten().reshape(-1,1)
out_w = w_after9[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(3, 2, 1)
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(3, 2, 2)
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

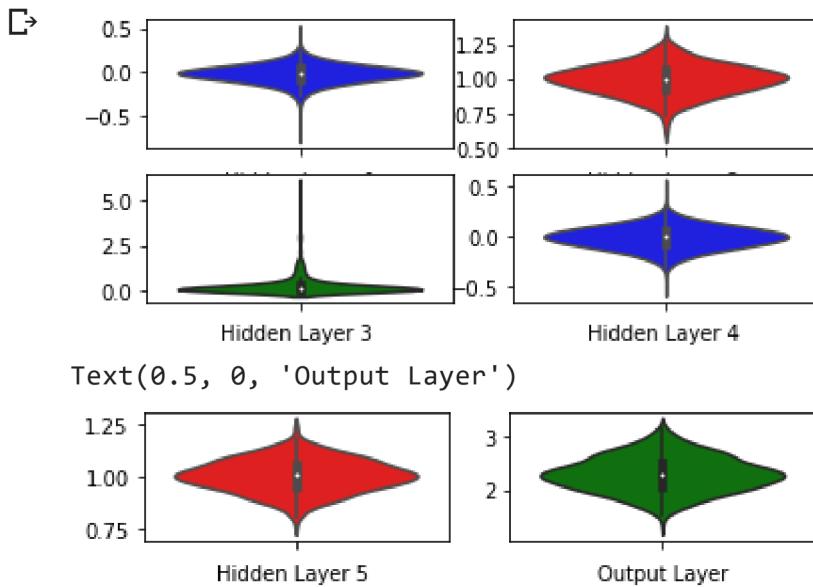
plt.subplot(3, 2, 3)
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(3, 2, 4)
ax = sns.violinplot(y=h4_w,color='b')
plt.xlabel('Hidden Layer 4')
plt.show()

```

```
plt.subplot(3, 2, 5)
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5')

plt.subplot(3, 2, 6)
ax = sns.violinplot(y=out_w, color='g')
plt.xlabel('Output Layer')
```



MLP + BN + Dropout + AdamOptimizer + ReLu activation

Part 3 : Here we use 5 Hidden Layers

Architecture 10: 784-512-456-392-256-128-10;
dropout: .5, .5

std = $\sqrt{2/\text{fan-in}}$ as it is ReLu

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
import warnings
warnings.filterwarnings('ignore')

model10 = Sequential()

model10.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.0233, seed=42))
model10.add(BatchNormalization())
model10.add(Dropout(0.5))

model10.add(Dense(456, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.06622, seed=42))
model10.add(BatchNormalization())
model10.add(Dropout(0.5))

model10.add(Dense(392, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.07142, seed=42))
model10.add(BatchNormalization())
model10.add(Dropout(0.5))

model10.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.08838, seed=42))
model10.add(BatchNormalization())
model10.add(Dropout(0.5))
```

```
model10.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=42))
model10.add(BatchNormalization())
model10.add(Dropout(0.5))

model10.add(Dense(output_dim, activation='softmax'))
model10.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model10.summary()

history10 = model10.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation
```

➡

Model: "sequential_23"

Layer (type)	Output Shape	Param #
dense_78 (Dense)	(None, 512)	401920
batch_normalization_56 (Batch Normalization)	(None, 512)	2048
dropout_56 (Dropout)	(None, 512)	0
dense_79 (Dense)	(None, 456)	233928
batch_normalization_57 (Batch Normalization)	(None, 456)	1824
dropout_57 (Dropout)	(None, 456)	0
dense_80 (Dense)	(None, 392)	179144
batch_normalization_58 (Batch Normalization)	(None, 392)	1568
dropout_58 (Dropout)	(None, 392)	0
dense_81 (Dense)	(None, 256)	100608
batch_normalization_59 (Batch Normalization)	(None, 256)	1024
dropout_59 (Dropout)	(None, 256)	0
dense_82 (Dense)	(None, 128)	32896
batch_normalization_60 (Batch Normalization)	(None, 128)	512
dropout_60 (Dropout)	(None, 128)	0
dense_83 (Dense)	(None, 10)	1290

Total params: 956,762

Trainable params: 953,274

Non-trainable params: 3,488

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 19s 314us/step - loss: 0.9767 - acc: 0.0000

Epoch 2/30

60000/60000 [=====] - 10s 173us/step - loss: 0.3401 - acc: 0.0000

Epoch 3/30

60000/60000 [=====] - 10s 173us/step - loss: 0.2488 - acc: 0.0000

Epoch 4/30

60000/60000 [=====] - 10s 175us/step - loss: 0.2093 - acc: 0.0000

Epoch 5/30

60000/60000 [=====] - 10s 174us/step - loss: 0.1804 - acc: 0.0000

Epoch 6/30

60000/60000 [=====] - 10s 174us/step - loss: 0.1651 - acc: 0.0000

Epoch 7/30

60000/60000 [=====] - 11s 177us/step - loss: 0.1526 - acc: 0.0000

Epoch 8/30

60000/60000 [=====] - 10s 173us/step - loss: 0.1375 - acc: 0.0000

```

Epoch 9/30
60000/60000 [=====] - 10s 174us/step - loss: 0.1343 - acc: 0
Epoch 10/30
60000/60000 [=====] - 10s 173us/step - loss: 0.1224 - acc: 0
Epoch 11/30
60000/60000 [=====] - 10s 171us/step - loss: 0.1142 - acc: 0
Epoch 12/30
60000/60000 [=====] - 10s 171us/step - loss: 0.1098 - acc: 0
Epoch 13/30
60000/60000 [=====] - 11s 178us/step - loss: 0.1027 - acc: 0
Epoch 14/30
60000/60000 [=====] - 11s 176us/step - loss: 0.1032 - acc: 0
Epoch 15/30
60000/60000 [=====] - 11s 176us/step - loss: 0.0977 - acc: 0
Epoch 16/30
60000/60000 [=====] - 11s 176us/step - loss: 0.0920 - acc: 0
Epoch 17/30
60000/60000 [=====] - 11s 176us/step - loss: 0.0917 - acc: 0
Epoch 18/30
60000/60000 [=====] - 11s 176us/step - loss: 0.0869 - acc: 0
Epoch 19/30
60000/60000 [=====] - 10s 172us/step - loss: 0.0859 - acc: 0
Epoch 20/30
60000/60000 [=====] - 11s 176us/step - loss: 0.0813 - acc: 0
Epoch 21/30
60000/60000 [=====] - 11s 175us/step - loss: 0.0759 - acc: 0
Epoch 22/30
60000/60000 [=====] - 11s 178us/step - loss: 0.0754 - acc: 0
Epoch 23/30
60000/60000 [=====] - 11s 181us/step - loss: 0.0711 - acc: 0
Epoch 24/30
60000/60000 [=====] - 11s 179us/step - loss: 0.0722 - acc: 0
Epoch 25/30
60000/60000 [=====] - 11s 176us/step - loss: 0.0696 - acc: 0
Epoch 26/30
60000/60000 [=====] - 11s 177us/step - loss: 0.0698 - acc: 0
Epoch 27/30
60000/60000 [=====] - 11s 177us/step - loss: 0.0648 - acc: 0
Epoch 28/30
60000/60000 [=====] - 11s 178us/step - loss: 0.0626 - acc: 0
Epoch 29/30
60000/60000 [=====] - 10s 174us/step - loss: 0.0596 - acc: 0
Epoch 30/30
60000/60000 [=====] - 11s 175us/step - loss: 0.0571 - acc: 0

```

```

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

score10 = model10.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score10[0])
print('Test accuracy:', score10[1])

```

```

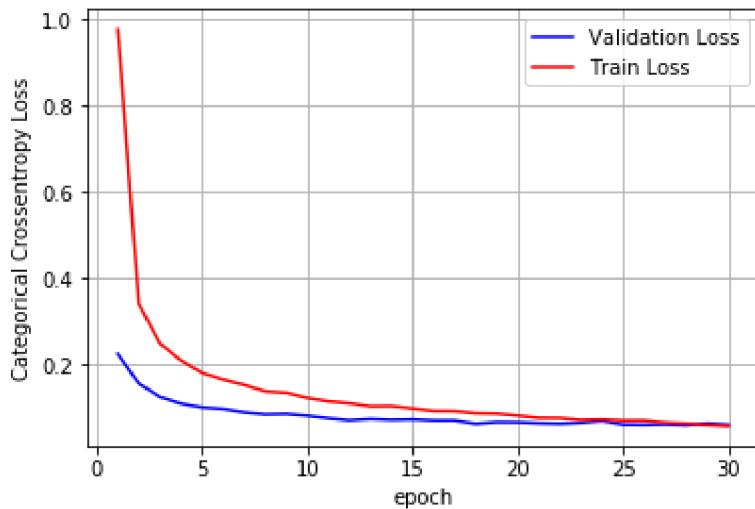
fig10,ax10 = plt.subplots(1,1)
ax10.set_xlabel('epoch') ; ax10.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x10 = list(range(1,nb_epoch+1))

vy10 = history10.history['val_loss']
ty10 = history10.history['loss']
plt_dynamic(x10, vy10, ty10, ax10)

```

↳ Test score: 0.06008605453777127
 Test accuracy: 0.9857



S.No	Architecture	Score	Test Accuracy
1.	784-366-100-10	0.0578912589	0.9844
2.	784-456-128-10	0.0567212487	0.984
3.	784-256-128-10	0.058956212	0.9823
4.	784-256-96-10	0.062314528	0.9832
—	—	—	—
5.	784-366-100-10	0.057912785	0.9846
6.	784-600-300-100-10	0.054455478	0.9857
7.	784-512-456-128-10	0.057914488	0.9839
8.	784-456-128-32-10	0.06451866	0.9845
—	—	—	—
9.	-600-500-400-200-128-	0.05667422	0.9854
10.	-512-456-392-256-128-	0.06008605	0.9857

▼ Model 1 ,2 ,3, 4 --> 2 Hidden Layers

Model 5 ,6 ,7, 8 --> 3 Hidden Layers

Model 9, 10 --> 5 Hidden Layers

So in general we do notice as numbers of layer increase Test Accuracy do increase however MNIST dataset is so simple that we can't be commenting on the basis of this. Also we see the dataset where first layer has lesser number of neurons compare to the last layer. Example whenever our first hidden layer is low like 256 or 366 (check model 3,4 and 5) accuracy is lower. Again we should not be "concluding" right away but in general these two points we do notice.